

Priority Inheritance with Backtracking for Iterative Multi-agent Path Finding

Keisuke Okumura¹, Manao Machida², Xaxier Défago¹, Yasumasa Tamura¹


1. Tokyo Institute of Technology, 2. NEC Corporation

SHORT DESCRIPTION

Propose an algorithm to plan paths for multiple moving agents iteratively/flexibly/scalably, called **PIBT**.

In biconnected-like graphs, PIBT ensures reachability, i.e., all agents reach their goals in finite time after being given.

MULTI-AGENT PATH FINDING (MAPF)



PRINCIPLES OF PIBT

In reality, MAPF must be solved **iteratively** requests to MAPF solvers

1. **Speed** over Optimality → Online, Real-time
2. **Decentralized** over Centralized → Robustness, Scalability
3. **Local** over Global Interaction → Concurrency, Scalability
4. Adaptivity for **Iterative** Use → Practicality


ALGORITHM

Prioritized Planning

pros: computationally-cheap
cons: **incomplete**

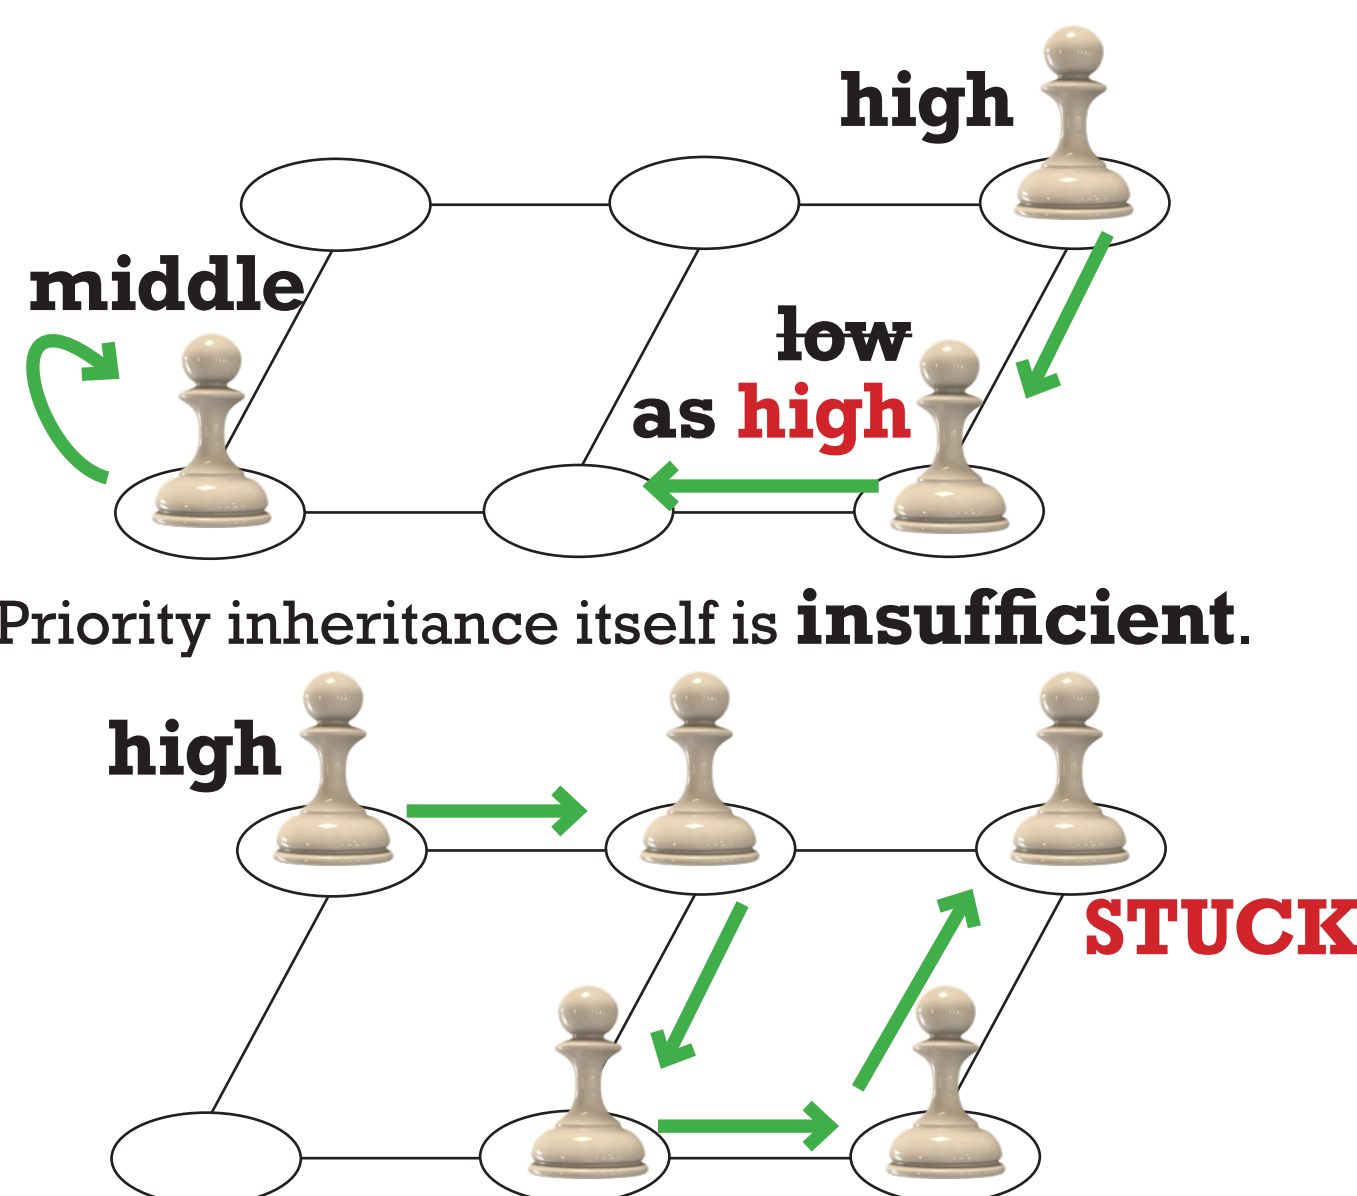
planning where time-window is just one

PIBT relies on



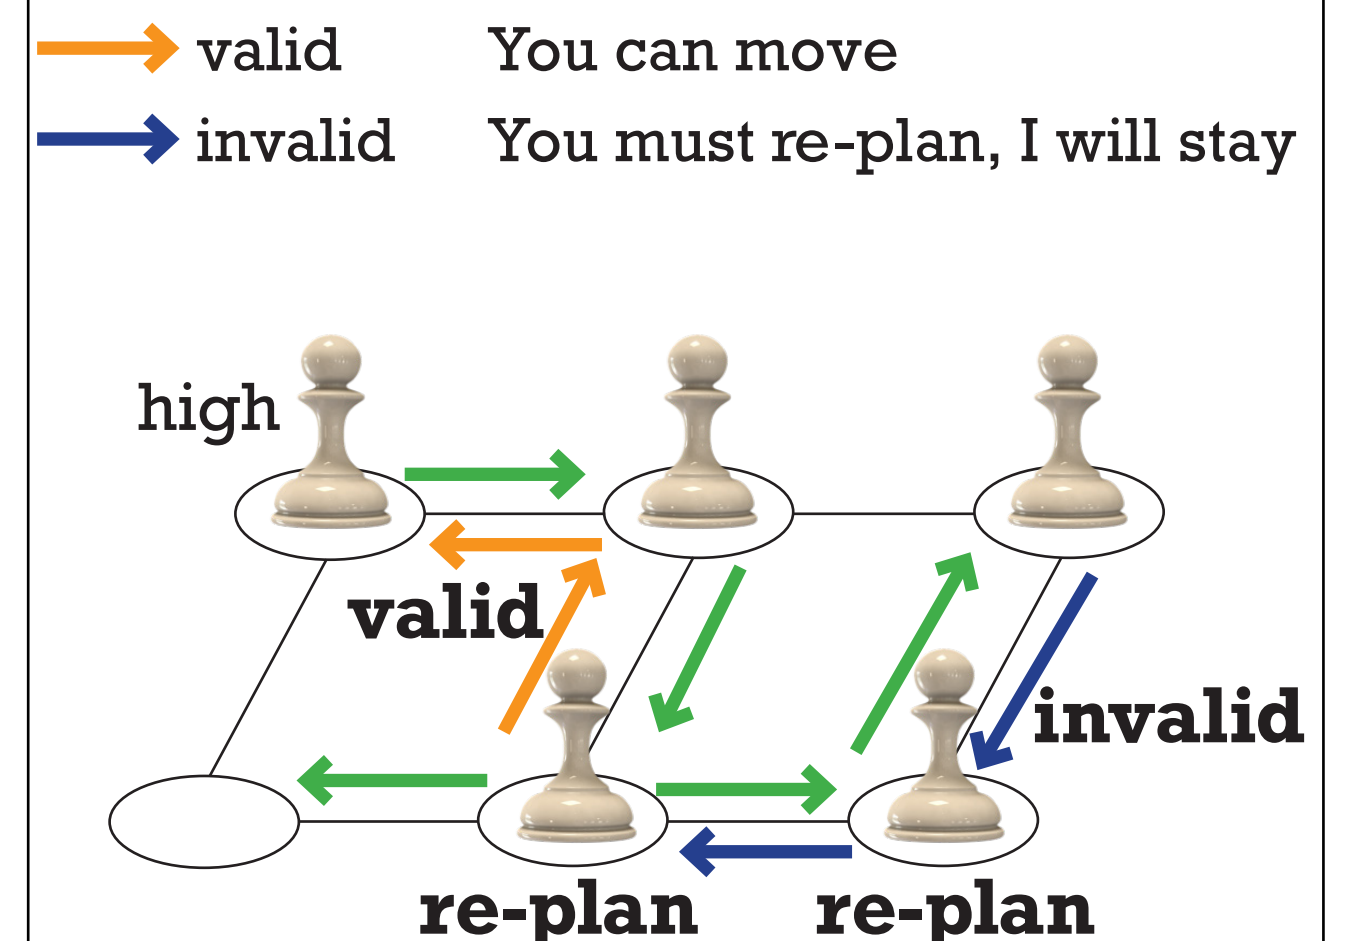
Priority Inheritance

The low-priority agent **temporarily inherits** the higher-priority.



Backtracking


Agents with priority inheritance have to wait for *backtracking*.



Dynamic Priorities

Once an agent reaches its goal, it drops its priority to ensure **fairness**.


increment priority (elapsed timestep)



Reachability


THEOREM 2 e.g.) biconnected graph

If a graph G has a simple cycle for all pairs of adjacent nodes, with PIBT, all agents reach their destination within $\text{diam}(G)|A|$ timesteps after being given.



Property

- computational complexity per timestep $O(|A| \cdot \Delta(G) \cdot F)$ maximum time required to choose the next target node
- locality



EVALUATION



Multi-agent Path Finding

PIBT


Parallel Push& Swap [Sajid et al., SoCS12]

WHCA* [Silver, AIIDE05]
window: 5(8x8), 10(otherwise)


- *path cost*: timestep when each agent reaches its goal and never move from then
- *makespan*: timestep when all agents reach their goals
- average over only success cases within the solver (in 100 instances) are shown




path cost




makespan



success



runtime



Multi-agent Pickup & Delivery

the problem for conveying packages in an automated warehouse [Ma et al., AAMAS17]

PIBT + task allocation
for each timestep, each free agent moves to the nearest pickup location of the non-assigned task

All tasks are completed in finite time if graph condition is satisfied.

50 agents, 500 tasks

PIBT TP [Ma et al., AAMAS17]

